

JOSHUA KAPPLER

San Francisco, CA • 512-210-9934 • Joshua.Kappler@gmail.com

github.com/JoshKappler • joshuakappler.com • linkedin.com/in/josh-kappler • youtube.com/boffy

Founding Developer Experience Engineer at General Translation. I build autonomous AI agents from scratch with direct SDK calls, no frameworks like LangChain, and I grew a YouTube channel to 2.1M subscribers and 270M views, so I can ship production systems and explain them in a way developers actually want to watch and read.

EXPERIENCE

Founding Developer Experience Engineer

2026 - Present

General Translation (full-stack localization infrastructure), San Francisco

- Owned the developer-facing surface end to end: the open-source **gt-next** and **gt-react** i18n libraries, the **gt** CLI, the docs search engine, and the translation platform dashboard.
- Built **Ask AI**, the dashboard's embedded assistant: it answers from live product docs through its own retrieval tools (search, fetch, per-page guides) and links users to the right org-scoped page.
- Designed and shipped the company's automated outbound growth engine: **multi-provider lead discovery and scoring** plus a **seven-tier email outreach sequencing system**.
- Ran a **security audit** of the platform and built the fixes: **cross-org permission leakage**, SSO/SCIM endpoints missing tenant checks, an AI-proxy allowlist and billing bypass, and prototype pollution via uploads.

Independent AI Engineer

2024 - 2026

Self-employed, San Francisco

- Shipped autonomous AI agent systems end to end for clients, including a deal-analysis platform in production at a private credit investment firm and a lead pipeline for a real-estate cash buyer. Direct SDK calls over framework abstractions, state machines for entity lifecycles, learning loops that record outcomes.
- Built multi-provider LLM orchestration: Anthropic API and Claude Agent SDK, OpenAI SDK, Groq, OpenRouter, with fallback routing, prompt caching on stable system prompts, and tool_use / function calling for structured output.
- Monorepo architecture with npm/pnpm workspaces, tRPC APIs, Drizzle + SQLite/Postgres, real-time dashboards. Validation at boundaries with Zod (TS) and Pydantic (Python).
- Built **fleetview**, a terminal cockpit that runs Claude Code agent fleets and keeps every machine in sync.

Founder & Content Creator, YouTube (Boffy)

2019 - Present

Independent digital media brand, 2.1M subscribers, 270M+ lifetime views

- Grew a channel from zero to 2.1M subscribers and 270M+ views across 136 long-form videos, mostly gaming.
- Wrote and presented every video on camera, including explainers that broke difficult technical processes down for a mass audience, like game modding and how graphics cards affect game performance.
- Negotiated brand partnerships with **RedMagic**, **Wargaming**, **GeoGuessr**, **YouTooZ**, **Factor**, and others. Represented products on camera and delivered scripted integrations on deadline.
- Hired and directed a team of 3 through the full content lifecycle, and ran a branded e-commerce store and a commercial Minecraft server.

PROJECTS

- memo-engine** (client build under NDA; live public demo): AI deal-analysis and investment-memo platform for a private credit firm. Runs **contextual retrieval** over pgvector with **ephemeral prompt caching**, producing institutional memos with **every claim cited back to source**. Durable orchestration via **Vercel Workflow DevKit**, forced tool_use + Zod for ~40-field extraction, reasoning on Claude Fable 5. Next.js 16, Neon Postgres, Anthropic SDK, Voyage AI.
- claim-wright** (built in a 36-hour weekend): Security-deposit insurance claim adjudication. **Claude Opus 4.8 reads** the lease, ledger, and move-out documents; a **pure-Python deterministic engine** applies the caps and rules, so a payout is never a number the model invented. **91% of decisions within \$250 of the human**, median error \$0, ~\$0.33 per claim. Forced tool_use, hybrid text/vision PDF reading, Django-Ninja, React.
- autohack**: Autonomous vulnerability hunter across HackerOne, Immunefi, Huntr, Bugcrowd, and Intigriti. **12-state findings pipeline** with 60-minute hunts, **adversarial self-review** before submission, and cross-hunt memory for learning. The same harness runs an Algora bounty agent that **spawns Claude Code sessions, runs tests, and opens PRs**. TypeScript monorepo, SQLite + Drizzle, PTY-based Claude spawning, tRPC.
- pinch**: Drive a real Claude Code session from an **Apple Watch** over cellular. Native watchOS SwiftUI client and a Node/TS backend running the **Claude Agent SDK** against live repos. **Durable at-least-once delivery** (persisted outbox, idempotent dedup) and **session resume** that rebuilds full context after a restart. HTTP short-poll transport, since watchOS blocks WebSockets on the watch path.

SKILLS

- AI / LLM**: Anthropic SDK (API + Claude Agent SDK), OpenAI SDK, Gemini, Groq, OpenRouter, Ollama, tool_use / function calling, prompt caching, multi-provider routing, agentic RAG
- Languages & Frameworks**: TypeScript, Python, JavaScript, Next.js 16, React 19, FastAPI, tRPC, Node.js, Tailwind CSS
- Data, Testing & Tooling**: PostgreSQL, SQLite, pgvector, Drizzle ORM, Zod, Pydantic, Playwright, Vitest, pytest, Claude Code, Git
- Developer Relations & Content**: On-camera presentation and video production, audience growth from zero to 2.1M, scriptwriting and project documentation, community engagement, brand and partner relationships

EDUCATION

- University of Oregon**: Computer Science (AI/ML specialization); left before finishing to build AI agent systems full-time.